

## Group Project: Cache Simulator

**Group Information Due:** Nov 2nd, 2018, Friday, 11:59pm

**Final Project Due:** Nov 30th, 2018, Friday, 11:59pm

### 1 Objectives

In this project, **you are to implement a cache simulator that simulates set associative caches with LRU replacement policy**. The goal of this project is to help you understand the internal operations of CPU caches and learn the typical procedure of computer architecture design.

### 2 General Project Descriptions

#### 2.1 Groups

Since this project requires implementation, testing, experimentation and documentation, you are required to work in groups. A group may have 2-3 students. Please submit a text file with the names of your group members to the Blackboard by **the end of Nov 2nd, Friday, 2018**. You should also signup for your group in the Blackboard by **the end of Nov 2nd, Friday, 2018**. To signup, please goto BlackBoard Tools  $\Rightarrow$  Groups. If you want to work individually on this project, please contact your instructor first.

#### 2.2 Programming Languages and Reference Systems

You are allowed to use the following programming languages: Python, C/C++ and Java. We strongly recommend that you to use Python2 with standard packages, which is available on every platform and will considerably simplify the implementation and submission process. A typical Python implementation does not exceed 500 lines of code. If you donot know how to program in Python, this is a good time to start learning it. As you are allowed to use three different languages, you will also need to submit a wrapper bash script to drive your simulator. More details about this wrapper script will be given later.

Our reference systems for this project are the CS department's fox servers. It can be accessed through ssh to addresses *fox01.cs.utsa.edu*, *fox02.cs.utsa.edu*, ..., *fox06.cs.utsa.edu*. We will grade your submissions on these servers. Therefore, please make sure your simulator can be correctly compiled and executed on these servers. Every student should have an account on these servers. If you have not changed your password before, your user name should be your *abcdid*, and your password should be your *banner id* (probably with the "@").

#### 2.3 Simulator Input and Memory Trace Files

Your simulator will need to take one command-line parameter that specifies the the path and name to a memory trace file. This memory trace file includes the memory addresses of a sequence of memory accesses. Your simulator should read in these memory addresses one by one and simulate their accesses to your cache.

More specially, A trace file includes lines in the following format:

```
#PC address:  #R/W #Memory Addresses
0x7f53a8cb3c23: W 0x7fffa80977a8
0x7f53a8cb4860: W 0x7fffa80977a0
0x7f53a8cb4864: W 0x7fffa8097798
... ..
```

The first column of each row is the PC (program counter or instruction address) of the instruction that generates this memory access. The second column indicates whether the access is a read or a write. The last column is the virtual memory address requested by this memory access. All columns are separated with a white space. The size of each memory access is not specified here, although you can assume there are no across-cacheline accesses. Moreover, although real caches typically use physical memory addresses as tags and to compute placement, for this project, you can just use these virtual memory addresses instead.

We provide several synthetic or real memory traces. You can download these traces from [here](#).<sup>1</sup>

Table 1 gives the detailed information for these traces.

Trace File	Description	Reference Cache Miss Rate
1KB_64B	Synthetic trace of accessing 1KB memory with 64B stride, repeated for two iterations	50% on 1KB, 16-way cache
4MB_4B	Synthetic trace of accessing 4MB memory with 4B stride, repeated for three iterations	2.08% on 4MB, 16-way cache
32MB_4B	Synthetic trace of accessing 32MB memory with 4B stride, repeated for three iterations	6.25% on 4MB, 16-way cache
bw_mem	Partial memory trace of benchmark <i>bw_mem</i>	1.56% on 4MB, 16-way cache
ls	Partial memory trace of an <i>ls</i> invocation	2.17% on 32KB, 16-way cache
gcc	Partial memory trace of a <i>gcc</i> compilation	1.89% on 32KB, 16-way cache
native_dgemm	Partial memory trace of a naive matrix multiplication of two 256x256 matrices with double-precision floats	50.24% on 256KB, 16-assoc cache
native_dgemm_full	Full memory trace of a naive matrix multiplication of two 256x256 matrices with double-precision floats	49.37% on 256KB, 16-assoc cache
openblas_dgemm	Partial memory trace of a cache-optimized matrix multiplication of two 512x512 matrices double-precision floats	8.30% on 256KB, 16-assoc cache
openblas_dgemm_full	Full memory trace of a cache-optimized matrix multiplication of two 512x512 matrices with double-precision floats	7.5% on 256KB, 16-assoc cache

Table 1: Memory trace files

The memory trace files are provided as *gzip* files. Please unzip them and use them directly as text files (although Python can read *gzip* files directly, please still unzip them). Your implementation should have the same cache miss rates as the reference miss rates for these traces. The reference miss rates are acquired with a popular cache simulator, *CPM\$im*.

The real memory traces are obtained with Intel *Pin* and Pintool *pinatrace*. For your experiments, you can download Intel Pin and the pintool to generate traces for other applications. You can also write some synthetic memory traces to test the correctness of your simulator.

## 2.4 Simulator Internals

Your simulator should be able to simulate set-associative caches of any sizes and any associativities/ways. It is a good idea to make the size and the way of your cache configurable through command-line parameters or configuration files. Note that, broadly speaking, direct mapped cache and fully-associative caches are also set-associative caches. Therefore, your simulator should be able to simulate these two types of cache as well.

<sup>1</sup>Full url is,  
[https://utsacloud-my.sharepoint.com/:f:/g/personal/wei\\_wang\\_utsa\\_edu/EtYXsykz0yxBguZl6K0Q7xsBTwQXzRf98PmurGhGKmihFQ?e=RsBi89](https://utsacloud-my.sharepoint.com/:f:/g/personal/wei_wang_utsa_edu/EtYXsykz0yxBguZl6K0Q7xsBTwQXzRf98PmurGhGKmihFQ?e=RsBi89)

For simulation, read in one line of the memory trace file, extract the memory address from the text, and simulate the access with the extract memory address to the cache. If the memory access misses the cache, you can assume that it is served by the memory.

Your simulator should implement at least the Least-Recently-Use (LRU) cache replacement policy. For the LRU policy, please at least implement a **true LRU** policy instead of the approximation algorithms learned in class. For a true LRU policy, you need to either use a queue or an extra field in the cache frame to record the access time stamp to each cache line. For write policy, please at least implement the **write-back policy**.

Your cache simulator is only required to simulate the standard operations of a cache. You do not need to simulate every physical function units used in the cache. Additionally, you probably want to read in one line at a time from the input memory trace file, instead of reading the whole files into the memory before processing it.

## 2.5 Simulator Outputs

Your simulator should output the simulation results to the screen (*standard out*, or *stdout*). The output should be a line that reads,

```
Cache miss rate: 66.67%
```

It is important that your simulator outputs only this line, so that we can grade all submissions using scripts. You are free to include other outputs during development phase to assist debugging or experimentation.

## 3 Project Procedure

In general, your project should carry out in the following steps.

1. **Forming a group.** Find teammates to form a group for the project. As stated previous, group information should be submitted in a text file to Blackboard by the end of Nov 2nd, 2018.
2. **Implementation.** Implement the cache simulator using your preferred language.
3. **Testing.** Test your simulator with different memory traces. First, you should write short memory traces to test different policies of your simulators, including the cache line identification policy, placement policy and replacement policy. After each individual policy is tested, you should test your simulator with large but regular memory traces (such as the provided synthetic memory traces) to ensure the simulator can handle accesses to large sizes of data. In the end, test with real memory traces to ensure your simulator can handle irregular memory accesses.
4. **Experimentation.** As we have seen in the class, no one cache size and associativity work the best for every application. Therefore, computer architects have to conduct various simulations to determine the cache size and associativity to include in the actual CPUs. In these simulations, they vary the cache sizes and the ways of their cache simulators, and simulate the memory accesses from memory traces obtained with real applications to determine the best configuration on average.

For this project, you will conduct similar experiments. More specially, you will take one real memory trace and simulate it on caches with different sizes and ways to determine which configuration gives the lowest cache miss rate. Repeat these experiments with at least one more real memory trace. Furthermore, besides the cache sizes, ways and provided traces, each group should at least evaluate another factor that affects cache performance, such as a none-LRU replacement policy, a none-writeback write policy or another real application's memory trace.

5. **Report.** At last, write a report on your simulator and experiments. Your report should have the following parts:
  - A cover page with group members' names, class name, report title and date

- A brief description of the implementation of the simulator.
- A description of your experiments and the experiment results with figures and/or tables.
- Individual contribution summary.

## 4 Deliverables and Submission Guidelines

Besides the simulator, please also include a wrapper bash script with the name “*run\_sim.sh*”, as discussed in Section 2.2. This bash script should accept one command line parameter, which is the path to a memory trace file. The script should then execute the simulator with the specified memory trace file. The simulator, in turn, simulates the memory accesses in the memory trace and outputs the cache miss rate. We only use this wrapper bash script to grade your simulator. The following commands and outputs show the expected behavior of this wrapper script.

```
$> ./run_sim.sh a_mem_trace.txt
Cache miss rate: 66.67%
```

If you choose to implement the simulator with C/C++, please also include a *makefile* in your zip to allow us to compile your source files. If you choose to implement the simulator with Java, please include just the compiled Java class files and write the Java command in the wrapper bash script to run these classes. If you choose to use Python3, please make sure your bash script uses *python3* as the python interpreter. Again, we strongly recommend every group use python2 with standard packages so that your simulator can be graded properly and simplify the development/submission process. For all submissions, especially for C/C++, Java and Python3 submissions, please make sure your code/script do compile and run correctly on fox servers. Otherwise, I am sure your code will not compile or run, and this will reflect negatively on your grades.

For the report, please submit it in one PDF file under the name “*report.pdf*”. The simulator should be configured with an LRU write-back cache of 1 MB with 16-ways.

Please zip your simulator files, the wrapper script and the report into one zip file, and submit the zip file to Blackboard. Please name the zip file “*cache\_sim.zip*”. There should be no directories in your zip file – all files should be at the root level.

Please strictly follow these submission guidelines. In the cases where the code/script fails to compile or execute, or the files are incorrectly named, or your simulator produces outputs in wrong format, there will be a 20% penalty if we can fix these errors. If we cannot fix these errors, the simulator will be considered as wrong and receives 0 points. All required file names are case-sensitive.

## 5 Grading

We will grade both the simulator and the report. For the simulator, we will execute it with several small memory traces to test if it can produce the correct cache miss rates. The memory trace files used in grading have the exact same format as the provided trace files. The report will be graded mainly based on the quality of implementation, the complexity and thoroughness of the experiments, and the quality of the writing. Individual grade will be negatively affected if a student does not exhibit a fair share of contribution.